

## Week 6 - Develop First modeling approach

```
In [1]: import pandas as pd
import numpy as np
# import dask.dataframe as dd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from datetime import datetime
from sklearn.metrics import roc_auc_score, f1_score, confusion_matrix
from sklearn.linear_model import LogisticRegression
from collections import defaultdict
```

```
In [2]: pd.set_option('display.max_columns', None)

train = pd.read_csv('train_final.csv', low_memory=False)
validation = pd.read_csv('val_set_final.csv')
```

```
In [3]: train.head()
```

```
Out[3]:
```

	Unnamed: 0	date	customer_code	employee_index	country	female	age	new_cust	seniority_
0	0	2015-07-28	664160	N	1	0	0.632653	0	
1	1	2016-01-28	1076784	N	1	0	0.214286	0	
2	2	2015-12-28	672465	N	1	0	0.387755	0	
3	3	2015-10-28	774528	N	1	0	0.397959	0	
4	4	2016-05-28	569598	N	1	0	0.459184	0	

```
In [4]: validation.head()
```

Out[4]:

	Unnamed: 0.1	Unnamed: 0	date	customer_code	employee_index	country_spain	female	age	fi
0	0	0	2015-11-28	161428	N	1	1	0.744898	
1	1	1	2015-12-28	367478	N	1	1	0.418367	
2	2	2	2015-11-28	643150	N	1	0	0.520408	
3	3	3	2016-04-28	1385854	N	1	0	0.367347	
4	4	4	2015-08-28	495733	N	1	0	0.346939	

Changing columns name and dropping columns so both datasets are the same

```
In [5]: train = train.rename(columns={'country': 'country_spain'})
```

```
In [6]: train = train.drop(columns=['Unnamed: 0'])
validation = validation.drop(columns=['Unnamed: 0', 'Unnamed: 0.1'])
drop = ['join_channel', 'province_name', 'employee_index', 'segment', 'total_products']
train = train.drop(columns=drop + ['customer_code_encoded'])
validation = validation.drop(columns=drop + ['payroll_acct.1', 'first_contract_date',
```

## Reading into the data

Setting products we want to predict

```
In [7]: products = ['savings_acct', 'guarantees', 'current_acct',
                    'derivada_acct', 'payroll_acct', 'junior_acct', 'mas_particular_acct',
                    'particular_acct', 'particular_plus_acct', 'short_term_depo',
                    'medium_term_depo', 'long_term_depo', 'e_acct', 'funds', 'mortgage',
                    'pension', 'loans', 'taxes', 'credit_card', 'securities', 'home_acct',
                    'pensions_2', 'direct_debt']
```

Dropping duplicates on customer code column since the last instance will show all the products a client has

```
In [8]: train = train.drop_duplicates(subset=['customer_code'], keep='last')
validation = validation.drop_duplicates(subset=['customer_code'], keep='last')

# Removing customers from validation set that appear in training set
validation = validation[~validation['customer_code'].isin(train['customer_code'])]
```

```
In [9]: print(train.info())
print(validation.info())
```

```
<class 'pandas.core.frame.DataFrame'>
Index: 706866 entries, 39288 to 6579716
Data columns (total 42 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   date                                  706866 non-null object
1   customer_code                        706866 non-null int64
2   country_spain                       706866 non-null int64
3   female                              706866 non-null int64
4   age                                  706866 non-null float64
5   new_cust                            706866 non-null int64
6   seniority_in_months                 706866 non-null float64
7   cust_type                           706866 non-null int64
8   residency_spain                    706866 non-null int64
9   birth_spain                        706866 non-null int64
10  active_cust                         706866 non-null int64
11  income                              706866 non-null float64
12  savings_acct                       706866 non-null int64
13  guarantees                         706866 non-null int64
14  current_acct                       706866 non-null int64
15  derivada_acct                      706866 non-null int64
16  payroll_acct                       706866 non-null int64
17  junior_acct                        706866 non-null int64
18  mas_particular_acct                706866 non-null int64
19  particular_acct                    706866 non-null int64
20  particular_plus_acct                706866 non-null int64
21  short_term_depo                    706866 non-null int64
22  medium_term_depo                   706866 non-null int64
23  long_term_depo                     706866 non-null int64
24  e_acct                             706866 non-null int64
25  funds                              706866 non-null int64
26  mortgage                           706866 non-null int64
27  pension                            706866 non-null int64
28  loans                              706866 non-null int64
29  taxes                              706866 non-null int64
30  credit_card                        706866 non-null int64
31  securities                         706866 non-null int64
32  home_acct                          706866 non-null int64
33  pensions_2                         706866 non-null int64
34  direct_debt                        706866 non-null int64
35  01 - TOP                           706866 non-null int64
36  02 - PARTICULARES                  706866 non-null int64
37  03 - UNIVERSITARIO                 706866 non-null int64
38  join_channel_encoded                706866 non-null float64
39  province_name_encoded               706866 non-null float64
40  employee_index_encoded              706866 non-null float64
41  income_to_age                       706866 non-null float64
```

dtypes: float64(7), int64(34), object(1)

memory usage: 231.9+ MB

None

```
<class 'pandas.core.frame.DataFrame'>
Index: 200333 entries, 51 to 2100200
Data columns (total 42 columns):
#   Column                                Non-Null Count  Dtype
---  -
0   date                                  200333 non-null object
1   customer_code                        200333 non-null int64
2   country_spain                       200333 non-null int64
3   female                              200333 non-null int64
4   age                                  200333 non-null float64
```

```

5  new_cust                200333 non-null  int64
6  seniority_in_months    200333 non-null  float64
7  cust_type              200333 non-null  int64
8  residency_spain        200333 non-null  int64
9  birth_spain            200333 non-null  int64
10 active_cust            200333 non-null  int64
11 income                 200333 non-null  float64
12 savings_acct          200333 non-null  int64
13 guarantees            200333 non-null  int64
14 current_acct          200333 non-null  int64
15 derivada_acct         200333 non-null  int64
16 payroll_acct          200333 non-null  int64
17 junior_acct           200333 non-null  int64
18 mas_particular_acct   200333 non-null  int64
19 particular_acct       200333 non-null  int64
20 particular_plus_acct  200333 non-null  int64
21 short_term_depo       200333 non-null  int64
22 medium_term_depo      200333 non-null  int64
23 long_term_depo        200333 non-null  int64
24 e_acct                200333 non-null  int64
25 funds                 200333 non-null  int64
26 mortgage              200333 non-null  int64
27 pension               200333 non-null  int64
28 loans                 200333 non-null  int64
29 taxes                 200333 non-null  int64
30 credit_card           200333 non-null  int64
31 securities            200333 non-null  int64
32 home_acct             200333 non-null  int64
33 pensions_2           200333 non-null  int64
34 direct_debt           200333 non-null  int64
35 01 - TOP              200333 non-null  int64
36 02 - PARTICULARES     200333 non-null  int64
37 03 - UNIVERSITARIO    200333 non-null  int64
38 join_channel_encoded  200333 non-null  float64
39 province_name_encoded 200333 non-null  float64
40 employee_index_encoded 200333 non-null  float64
41 income_to_age         200333 non-null  float64

```

dtypes: float64(7), int64(34), object(1)

memory usage: 65.7+ MB

None

<class 'pandas.core.frame.DataFrame'>

Index: 200333 entries, 51 to 2100200

Data columns (total 42 columns):

#	Column	Non-Null Count	Dtype
0	date	200333 non-null	object
1	customer_code	200333 non-null	int64
2	country_spain	200333 non-null	int64
3	female	200333 non-null	int64
4	age	200333 non-null	float64
5	new_cust	200333 non-null	int64
6	seniority_in_months	200333 non-null	float64
7	cust_type	200333 non-null	int64
8	residency_spain	200333 non-null	int64
9	birth_spain	200333 non-null	int64
10	active_cust	200333 non-null	int64
11	income	200333 non-null	float64
12	savings_acct	200333 non-null	int64
13	guarantees	200333 non-null	int64
14	current_acct	200333 non-null	int64

```

15  derivada_acct          200333 non-null  int64
16  payroll_acct          200333 non-null  int64
17  junior_acct           200333 non-null  int64
18  mas_particular_acct   200333 non-null  int64
19  particular_acct       200333 non-null  int64
20  particular_plus_acct  200333 non-null  int64
21  short_term_depo       200333 non-null  int64
22  medium_term_depo      200333 non-null  int64
23  long_term_depo        200333 non-null  int64
24  e_acct                200333 non-null  int64
25  funds                 200333 non-null  int64
26  mortgage              200333 non-null  int64
27  pension               200333 non-null  int64
28  loans                 200333 non-null  int64
29  taxes                 200333 non-null  int64
30  credit_card           200333 non-null  int64
31  securities            200333 non-null  int64
32  home_acct             200333 non-null  int64
33  pensions_2            200333 non-null  int64
34  direct_debt           200333 non-null  int64
35  01 - TOP              200333 non-null  int64
36  02 - PARTICULARES     200333 non-null  int64
37  03 - UNIVERSITARIO    200333 non-null  int64
38  join_channel_encoded   200333 non-null  float64
39  province_name_encoded  200333 non-null  float64
40  employee_index_encoded 200333 non-null  float64
41  income_to_age          200333 non-null  float64
dtypes: float64(7), int64(34), object(1)
memory usage: 65.7+ MB
None

```

## Pre-processing

Defining our Xs and Ys

```

In [10]: X_train = train.drop(['customer_code', 'date'] + products, axis=1)
         y_train = train[products]

         X_val = validation.drop(['customer_code', 'date'] + products, axis=1)
         y_val = validation[products]

```

```

In [11]: print("Shape of X_train:", X_train.shape)
         print("Shape of y_train:", y_train.shape)

         print("Shape of X_val:", X_val.shape)
         print("Shape of y_val:", y_val.shape)

```

```

Shape of X_train: (706866, 17)
Shape of y_train: (706866, 23)
Shape of X_val: (200333, 17)
Shape of y_val: (200333, 23)

```

## Training

```

In [12]: # Hyperparameters
         hyperparameter_variations = [

```

```

{'C': 0.01, 'solver': 'liblinear', 'max_iter': 100},
{'C': 1, 'solver': 'lbfgs', 'max_iter': 500},
{'C': 10, 'solver': 'liblinear', 'max_iter': 300},
]

```

```

In [13]: # Storing trained models and predictions
models = {}
metrics = defaultdict(lambda: defaultdict(dict))

```

We will create a model to train on the training data using all 3 hyperparameters we set. We will use this trained model to predict the product recommendations on the validation set and compare the results between the different hyperparameters and different metrics we chose to use, which are ROC AUC, F1 Score and Confusion Matrix.

We will calculate ROC AUC using probabilities (predict\_proba() method), which is more appropriate for this metric since ROC AUC works with predicted probabilities for the positive class and not binary predictions.

F1 Score and confusion matrix were calculated using the binary predictions (predict() method), which is the correct approach for these metrics.

```

In [14]: # Train and evaluate each hyperparameter variation
for i, params in enumerate(hyperparameter_variations):
    print(f"\nTraining variation {i + 1} with parameters: {params}")

    for product in products:
        print(f"Training model for product: {product}")
        clf = LogisticRegression(**params)

        # Target column for current product
        y_train_product = y_train[product].values
        y_val_product = y_val[product].values

        # Train the model on current product
        clf.fit(X_train, y_train_product)

        # Make predictions
        y_train_pred = clf.predict(X_train)
        y_val_pred = clf.predict(X_val)
        y_train_pred_proba = clf.predict_proba(X_train)[ :, 1]
        y_val_pred_proba = clf.predict_proba(X_val)[ :, 1]

        # Calculate metrics for training set and validation sets
        metrics[f'Variation {i + 1}']['train'][product] = {
            'ROC AUC': roc_auc_score(y_train_product, y_train_pred_proba),
            'F1 Score': f1_score(y_train_product, y_train_pred),
            'Confusion Matrix': confusion_matrix(y_train_product, y_train_pred)
        }

        metrics[f'Variation {i + 1}']['val'][product] = {
            'ROC AUC': roc_auc_score(y_val_product, y_val_pred_proba),
            'F1 Score': f1_score(y_val_product, y_val_pred),
            'Confusion Matrix': confusion_matrix(y_val_product, y_val_pred)
        }

    print(f"\nResults for '{product}' in variation {i + 1}:")
    print(f"Training - ROC AUC: {metrics[f'Variation {i + 1}']['train'][product]['ROC AUC']}")
    print(f"Training - F1 Score: {metrics[f'Variation {i + 1}']['train'][product]['F1 Score']}")

```

```
print(f"Validation - ROC AUC: {metrics[f'Variation {i + 1}']['val'][product]['  
      f"F1 Score: {metrics[f'Variation {i + 1}']['val'][product]['F1 Score']:.2f}
```

Training variation 1 with parameters: {'C': 0.01, 'solver': 'liblinear', 'max\_iter': 100}

Training model for product: savings\_acct

Results for 'savings\_acct' in variation 1:

Training - ROC AUC: 0.4029, F1 Score: 0.0000

Validation - ROC AUC: 0.4489, F1 Score: 0.0000

Training model for product: guarantees

Results for 'guarantees' in variation 1:

Training - ROC AUC: 0.2619, F1 Score: 0.0000

Validation - ROC AUC: 0.1172, F1 Score: 0.0000

Training model for product: current\_acct

Results for 'current\_acct' in variation 1:

Training - ROC AUC: 0.7485, F1 Score: 0.7865

Validation - ROC AUC: 0.7564, F1 Score: 0.7899

Training model for product: derivada\_acct

Results for 'derivada\_acct' in variation 1:

Training - ROC AUC: 0.7772, F1 Score: 0.0000

Validation - ROC AUC: 0.7503, F1 Score: 0.0000

Training model for product: payroll\_acct

Results for 'payroll\_acct' in variation 1:

Training - ROC AUC: 0.8646, F1 Score: 0.0001

Validation - ROC AUC: 0.8595, F1 Score: 0.0075

Training model for product: junior\_acct

Results for 'junior\_acct' in variation 1:

Training - ROC AUC: 0.9990, F1 Score: 0.1412

Validation - ROC AUC: 0.9972, F1 Score: 0.1272

Training model for product: mas\_particular\_acct

Results for 'mas\_particular\_acct' in variation 1:

Training - ROC AUC: 0.8256, F1 Score: 0.0000

Validation - ROC AUC: 0.8864, F1 Score: 0.0000

Training model for product: particular\_acct

Results for 'particular\_acct' in variation 1:

Training - ROC AUC: 0.8847, F1 Score: 0.2141

Validation - ROC AUC: 0.9270, F1 Score: 0.2457

Training model for product: particular\_plus\_acct

Results for 'particular\_plus\_acct' in variation 1:

Training - ROC AUC: 0.8132, F1 Score: 0.0000

Validation - ROC AUC: 0.8590, F1 Score: 0.0288

Training model for product: short\_term\_depo

Results for 'short\_term\_depo' in variation 1:

Training - ROC AUC: 0.9249, F1 Score: 0.0000

Validation - ROC AUC: 0.9126, F1 Score: 0.0000

Training model for product: medium\_term\_depo

Results for 'medium\_term\_depo' in variation 1:

Training - ROC AUC: 0.8772, F1 Score: 0.0000

Validation - ROC AUC: 0.8977, F1 Score: 0.0000

Training model for product: long\_term\_depo

Results for 'long\_term\_depo' in variation 1:



Training - ROC AUC: 0.9263, F1 Score: 0.3412  
Validation - ROC AUC: 0.9359, F1 Score: 0.2931  
Training model for product: e\_acct

Results for 'e\_acct' in variation 1:  
Training - ROC AUC: 0.8606, F1 Score: 0.2165  
Validation - ROC AUC: 0.8769, F1 Score: 0.1683  
Training model for product: funds

Results for 'funds' in variation 1:  
Training - ROC AUC: 0.9216, F1 Score: 0.0000  
Validation - ROC AUC: 0.9357, F1 Score: 0.0584  
Training model for product: mortgage

Results for 'mortgage' in variation 1:  
Training - ROC AUC: 0.9208, F1 Score: 0.0000  
Validation - ROC AUC: 0.9466, F1 Score: 0.0000  
Training model for product: pension

Results for 'pension' in variation 1:  
Training - ROC AUC: 0.9208, F1 Score: 0.0000  
Validation - ROC AUC: 0.9321, F1 Score: 0.0378  
Training model for product: loans

Results for 'loans' in variation 1:  
Training - ROC AUC: 0.8397, F1 Score: 0.0000  
Validation - ROC AUC: 0.8581, F1 Score: 0.0000  
Training model for product: taxes

Results for 'taxes' in variation 1:  
Training - ROC AUC: 0.8592, F1 Score: 0.0008  
Validation - ROC AUC: 0.8657, F1 Score: 0.0307  
Training model for product: credit\_card

Results for 'credit\_card' in variation 1:  
Training - ROC AUC: 0.8904, F1 Score: 0.0001  
Validation - ROC AUC: 0.9184, F1 Score: 0.0405  
Training model for product: securities

Results for 'securities' in variation 1:  
Training - ROC AUC: 0.9137, F1 Score: 0.0002  
Validation - ROC AUC: 0.9285, F1 Score: 0.0860  
Training model for product: home\_acct

Results for 'home\_acct' in variation 1:  
Training - ROC AUC: 0.8714, F1 Score: 0.0000  
Validation - ROC AUC: 0.8997, F1 Score: 0.0000  
Training model for product: pensions\_2

Results for 'pensions\_2' in variation 1:  
Training - ROC AUC: 0.8614, F1 Score: 0.0001  
Validation - ROC AUC: 0.8587, F1 Score: 0.0091  
Training model for product: direct\_debt

Results for 'direct\_debt' in variation 1:  
Training - ROC AUC: 0.8688, F1 Score: 0.0356  
Validation - ROC AUC: 0.8709, F1 Score: 0.0514

Training variation 2 with parameters: {'C': 1, 'solver': 'lbfgs', 'max\_iter': 500}  
Training model for product: savings\_acct

Results for 'savings\_acct' in variation 2:  
 Training - ROC AUC: 0.8689, F1 Score: 0.0000  
 Validation - ROC AUC: 0.9334, F1 Score: 0.0000  
 Training model for product: guarantees

Results for 'guarantees' in variation 2:  
 Training - ROC AUC: 0.9642, F1 Score: 0.0000  
 Validation - ROC AUC: 0.9887, F1 Score: 0.0000  
 Training model for product: current\_acct

c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
 n\_iter\_i = \_check\_optimize\_result(

Results for 'current\_acct' in variation 2:  
 Training - ROC AUC: 0.7480, F1 Score: 0.7863  
 Validation - ROC AUC: 0.7569, F1 Score: 0.7895  
 Training model for product: derivada\_acct

c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
 n\_iter\_i = \_check\_optimize\_result(

Results for 'derivada\_acct' in variation 2:  
 Training - ROC AUC: 0.8804, F1 Score: 0.0000  
 Validation - ROC AUC: 0.9080, F1 Score: 0.0000  
 Training model for product: payroll\_acct

c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
 n\_iter\_i = \_check\_optimize\_result(

Results for 'payroll\_acct' in variation 2:  
 Training - ROC AUC: 0.8663, F1 Score: 0.0003  
 Validation - ROC AUC: 0.8602, F1 Score: 0.0090  
 Training model for product: junior\_acct

c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
 STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
 n\_iter\_i = \_check\_optimize\_result(

Results for 'junior\_acct' in variation 2:  
Training - ROC AUC: 0.9994, F1 Score: 0.8811  
Validation - ROC AUC: 0.9986, F1 Score: 0.8270  
Training model for product: mas\_particular\_acct

Results for 'mas\_particular\_acct' in variation 2:  
Training - ROC AUC: 0.8412, F1 Score: 0.0000  
Validation - ROC AUC: 0.8851, F1 Score: 0.0000  
Training model for product: particular\_acct

Results for 'particular\_acct' in variation 2:  
Training - ROC AUC: 0.8848, F1 Score: 0.2291  
Validation - ROC AUC: 0.9254, F1 Score: 0.2534  
Training model for product: particular\_plus\_acct

c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
n\_iter\_i = \_check\_optimize\_result(

Results for 'particular\_plus\_acct' in variation 2:  
Training - ROC AUC: 0.8138, F1 Score: 0.0000  
Validation - ROC AUC: 0.8616, F1 Score: 0.0298  
Training model for product: short\_term\_depo

Results for 'short\_term\_depo' in variation 2:  
Training - ROC AUC: 0.9455, F1 Score: 0.0000  
Validation - ROC AUC: 0.9457, F1 Score: 0.0058  
Training model for product: medium\_term\_depo

Results for 'medium\_term\_depo' in variation 2:  
Training - ROC AUC: 0.8952, F1 Score: 0.0000  
Validation - ROC AUC: 0.9271, F1 Score: 0.0377  
Training model for product: long\_term\_depo

c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
n\_iter\_i = \_check\_optimize\_result(

Results for 'long\_term\_depo' in variation 2:  
Training - ROC AUC: 0.9270, F1 Score: 0.3508  
Validation - ROC AUC: 0.9384, F1 Score: 0.3035  
Training model for product: e\_acct

Results for 'e\_acct' in variation 2:  
Training - ROC AUC: 0.8608, F1 Score: 0.2218  
Validation - ROC AUC: 0.8738, F1 Score: 0.1717  
Training model for product: funds

```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Results for 'funds' in variation 2:

Training - ROC AUC: 0.9232, F1 Score: 0.0044

Validation - ROC AUC: 0.9381, F1 Score: 0.0611

Training model for product: mortgage

Results for 'mortgage' in variation 2:

Training - ROC AUC: 0.9267, F1 Score: 0.0000

Validation - ROC AUC: 0.9486, F1 Score: 0.0059

Training model for product: pension

```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Results for 'pension' in variation 2:

Training - ROC AUC: 0.9223, F1 Score: 0.0056

Validation - ROC AUC: 0.9338, F1 Score: 0.0317

Training model for product: loans

```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Results for 'loans' in variation 2:

Training - ROC AUC: 0.8544, F1 Score: 0.0000

Validation - ROC AUC: 0.8878, F1 Score: 0.0000

Training model for product: taxes

Results for 'taxes' in variation 2:

Training - ROC AUC: 0.8597, F1 Score: 0.0012

Validation - ROC AUC: 0.8660, F1 Score: 0.0311

Training model for product: credit\_card

```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
```

Results for 'credit\_card' in variation 2:

Training - ROC AUC: 0.8906, F1 Score: 0.0061

Validation - ROC AUC: 0.9202, F1 Score: 0.0440

Training model for product: securities

c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

Results for 'securities' in variation 2:

Training - ROC AUC: 0.9142, F1 Score: 0.0074

Validation - ROC AUC: 0.9301, F1 Score: 0.0965

Training model for product: home\_acct

c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

Results for 'home\_acct' in variation 2:

Training - ROC AUC: 0.8889, F1 Score: 0.0000

Validation - ROC AUC: 0.9137, F1 Score: 0.0091

Training model for product: pensions\_2

c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

Results for 'pensions\_2' in variation 2:

Training - ROC AUC: 0.8623, F1 Score: 0.0013

Validation - ROC AUC: 0.8596, F1 Score: 0.0105

Training model for product: direct\_debt

c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear\_model\\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

n\_iter\_i = \_check\_optimize\_result(

Results for 'direct\_debt' in variation 2:

Training - ROC AUC: 0.8691, F1 Score: 0.0500

Validation - ROC AUC: 0.8709, F1 Score: 0.0606

Training variation 3 with parameters: {'C': 10, 'solver': 'liblinear', 'max\_iter': 300}

Training model for product: savings\_acct

Results for 'savings\_acct' in variation 3:

Training - ROC AUC: 0.8729, F1 Score: 0.0000

Validation - ROC AUC: 0.9315, F1 Score: 0.0000

Training model for product: guarantees

Results for 'guarantees' in variation 3:

Training - ROC AUC: 0.9734, F1 Score: 0.0000

Validation - ROC AUC: 0.9106, F1 Score: 0.0000

Training model for product: current\_acct

Results for 'current\_acct' in variation 3:

Training - ROC AUC: 0.7490, F1 Score: 0.7866

Validation - ROC AUC: 0.7587, F1 Score: 0.7901

Training model for product: derivada\_acct

Results for 'derivada\_acct' in variation 3:

Training - ROC AUC: 0.8808, F1 Score: 0.0000

Validation - ROC AUC: 0.8979, F1 Score: 0.0000

Training model for product: payroll\_acct

Results for 'payroll\_acct' in variation 3:

Training - ROC AUC: 0.8663, F1 Score: 0.0010

Validation - ROC AUC: 0.8642, F1 Score: 0.0090

Training model for product: junior\_acct

Results for 'junior\_acct' in variation 3:

Training - ROC AUC: 0.9996, F1 Score: 0.8854

Validation - ROC AUC: 0.9996, F1 Score: 0.8401

Training model for product: mas\_particular\_acct

Results for 'mas\_particular\_acct' in variation 3:

Training - ROC AUC: 0.8412, F1 Score: 0.0000

Validation - ROC AUC: 0.8902, F1 Score: 0.0000

Training model for product: particular\_acct

Results for 'particular\_acct' in variation 3:

Training - ROC AUC: 0.8849, F1 Score: 0.2309

Validation - ROC AUC: 0.9303, F1 Score: 0.2574

Training model for product: particular\_plus\_acct

Results for 'particular\_plus\_acct' in variation 3:

Training - ROC AUC: 0.8137, F1 Score: 0.0000

Validation - ROC AUC: 0.8670, F1 Score: 0.0000

Training model for product: short\_term\_depo

Results for 'short\_term\_depo' in variation 3:

Training - ROC AUC: 0.9473, F1 Score: 0.0000

Validation - ROC AUC: 0.9438, F1 Score: 0.0126

Training model for product: medium\_term\_depo

Results for 'medium\_term\_depo' in variation 3:

Training - ROC AUC: 0.8952, F1 Score: 0.0000

Validation - ROC AUC: 0.9285, F1 Score: 0.0358  
Training model for product: long\_term\_depo

Results for 'long\_term\_depo' in variation 3:  
Training - ROC AUC: 0.9270, F1 Score: 0.3501  
Validation - ROC AUC: 0.9401, F1 Score: 0.3052  
Training model for product: e\_acct

Results for 'e\_acct' in variation 3:  
Training - ROC AUC: 0.8612, F1 Score: 0.2203  
Validation - ROC AUC: 0.8791, F1 Score: 0.1709  
Training model for product: funds

Results for 'funds' in variation 3:  
Training - ROC AUC: 0.9231, F1 Score: 0.0042  
Validation - ROC AUC: 0.9420, F1 Score: 0.0597  
Training model for product: mortgage

Results for 'mortgage' in variation 3:  
Training - ROC AUC: 0.9270, F1 Score: 0.0000  
Validation - ROC AUC: 0.9518, F1 Score: 0.0079  
Training model for product: pension

Results for 'pension' in variation 3:  
Training - ROC AUC: 0.9224, F1 Score: 0.0056  
Validation - ROC AUC: 0.9397, F1 Score: 0.0367  
Training model for product: loans

Results for 'loans' in variation 3:  
Training - ROC AUC: 0.8544, F1 Score: 0.0000  
Validation - ROC AUC: 0.8876, F1 Score: 0.0000  
Training model for product: taxes

Results for 'taxes' in variation 3:  
Training - ROC AUC: 0.8598, F1 Score: 0.0012  
Validation - ROC AUC: 0.8734, F1 Score: 0.0234  
Training model for product: credit\_card

Results for 'credit\_card' in variation 3:  
Training - ROC AUC: 0.8907, F1 Score: 0.0063  
Validation - ROC AUC: 0.9235, F1 Score: 0.0438  
Training model for product: securities

Results for 'securities' in variation 3:  
Training - ROC AUC: 0.9142, F1 Score: 0.0074  
Validation - ROC AUC: 0.9320, F1 Score: 0.0970  
Training model for product: home\_acct

Results for 'home\_acct' in variation 3:  
Training - ROC AUC: 0.8892, F1 Score: 0.0000  
Validation - ROC AUC: 0.9174, F1 Score: 0.0000  
Training model for product: pensions\_2

Results for 'pensions\_2' in variation 3:  
Training - ROC AUC: 0.8625, F1 Score: 0.0010  
Validation - ROC AUC: 0.8615, F1 Score: 0.0107  
Training model for product: direct\_debt

Results for 'direct\_debt' in variation 3:

Training - ROC AUC: 0.8693, F1 Score: 0.0468  
Validation - ROC AUC: 0.8705, F1 Score: 0.0586

Creating a table to see the results in a easier to interpret way

```
In [15]: train_metrics_df = pd.DataFrame.from_dict({(i,j): metrics[i]['train'][j]
            for i in metrics.keys()
            for j in products},
            orient='index')

val_metrics_df = pd.DataFrame.from_dict({(i,j): metrics[i]['val'][j]
            for i in metrics.keys()
            for j in products},
            orient='index')
```

```
In [16]: pd.set_option('display.max_rows', None)
print("Training Metrics Table:")
train_metrics_df
```

Training Metrics Table:



Out[16]:

		ROC AUC	F1 Score	Confusion Matrix
Variation 1	savings_acct	0.402916	0.000000	[[706792, 0], [74, 0]]
	guarantees	0.261854	0.000000	[[706850, 0], [16, 0]]
	current_acct	0.748522	0.786467	[[150022, 126938], [69026, 360880]]
	derivada_acct	0.777204	0.000000	[[706603, 0], [263, 0]]
	payroll_acct	0.864570	0.000103	[[667859, 2], [39003, 2]]
	junior_acct	0.998986	0.141183	[[700283, 2], [6081, 500]]
	mas_particular_acct	0.825569	0.000000	[[701121, 0], [5745, 0]]
	particular_acct	0.884666	0.214141	[[600002, 20131], [73919, 12814]]
	particular_plus_acct	0.813231	0.000000	[[677357, 1], [29508, 0]]
	short_term_depo	0.924857	0.000000	[[705857, 0], [1009, 0]]
	medium_term_depo	0.877228	0.000000	[[705817, 0], [1049, 0]]
	long_term_depo	0.926317	0.341200	[[672212, 5132], [22394, 7128]]
	e_acct	0.860581	0.216472	[[640624, 7361], [50841, 8040]]
	funds	0.921595	0.000000	[[694096, 3], [12767, 0]]
	mortgage	0.920838	0.000000	[[702807, 0], [4059, 0]]
	pension	0.920788	0.000000	[[700453, 1], [6412, 0]]
	loans	0.839712	0.000000	[[705210, 0], [1656, 0]]
	taxes	0.859193	0.000784	[[668627, 17], [38207, 15]]
	credit_card	0.890356	0.000129	[[675775, 3], [31086, 2]]
	securities	0.913740	0.000227	[[689223, 12], [17629, 2]]
Variation 2	home_acct	0.871426	0.000000	[[704156, 0], [2710, 0]]
	pensions_2	0.861401	0.000094	[[664258, 2], [42604, 2]]
	direct_debt	0.868809	0.035647	[[615556, 1308], [88345, 1657]]
	savings_acct	0.868903	0.000000	[[706792, 0], [74, 0]]
	guarantees	0.964197	0.000000	[[706850, 0], [16, 0]]
	current_acct	0.748034	0.786279	[[150207, 126753], [69288, 360618]]
	derivada_acct	0.880394	0.000000	[[706603, 0], [263, 0]]
	payroll_acct	0.866318	0.000256	[[667856, 5], [39000, 5]]
	junior_acct	0.999425	0.881061	[[699481, 804], [766, 5815]]
	mas_particular_acct	0.841209	0.000000	[[701121, 0], [5745, 0]]
	particular_acct	0.884820	0.229081	[[598646, 21487], [72734, 13999]]
	particular_plus_acct	0.813776	0.000000	[[677356, 2], [29508, 0]]
	short_term_depo	0.945518	0.000000	[[705857, 0], [1009, 0]]
	medium_term_depo	0.895236	0.000000	[[705817, 0], [1049, 0]]

		ROC AUC	F1 Score	Confusion Matrix
	<b>long_term_depo</b>	0.927018	0.350823	[[671929, 5415], [22090, 7432]]
	<b>e_acct</b>	0.860760	0.221820	[[640475, 7510], [50599, 8282]]
	<b>funds</b>	0.923179	0.004361	[[694052, 47], [12739, 28]]
	<b>mortgage</b>	0.926707	0.000000	[[702806, 1], [4059, 0]]
	<b>pension</b>	0.922319	0.005588	[[700442, 12], [6394, 18]]
	<b>loans</b>	0.854446	0.000000	[[705210, 0], [1656, 0]]
	<b>taxes</b>	0.859697	0.001202	[[668620, 24], [38199, 23]]
	<b>credit_card</b>	0.890593	0.006080	[[675713, 65], [30993, 95]]
	<b>securities</b>	0.914192	0.007435	[[689178, 57], [17565, 66]]
	<b>home_acct</b>	0.888943	0.000000	[[704156, 0], [2710, 0]]
	<b>pensions_2</b>	0.862302	0.001266	[[664248, 12], [42579, 27]]
<b>Variation 3</b>	<b>direct_debt</b>	0.869138	0.049997	[[614470, 2394], [87633, 2369]]
	<b>savings_acct</b>	0.872940	0.000000	[[706792, 0], [74, 0]]
	<b>guarantees</b>	0.973446	0.000000	[[706850, 0], [16, 0]]
	<b>current_acct</b>	0.748993	0.786626	[[149726, 127234], [68714, 361192]]
	<b>derivada_acct</b>	0.880811	0.000000	[[706603, 0], [263, 0]]
	<b>payroll_acct</b>	0.866259	0.001025	[[667849, 12], [38985, 20]]
	<b>junior_acct</b>	0.999570	0.885389	[[699470, 815], [706, 5875]]
	<b>mas_particular_acct</b>	0.841194	0.000000	[[701121, 0], [5745, 0]]
	<b>particular_acct</b>	0.884887	0.230861	[[598531, 21602], [72596, 14137]]
	<b>particular_plus_acct</b>	0.813682	0.000000	[[677358, 0], [29508, 0]]
	<b>short_term_depo</b>	0.947280	0.000000	[[705857, 0], [1009, 0]]
	<b>medium_term_depo</b>	0.895238	0.000000	[[705817, 0], [1049, 0]]
	<b>long_term_depo</b>	0.927018	0.350078	[[671952, 5392], [22114, 7408]]
	<b>e_acct</b>	0.861151	0.220272	[[640548, 7437], [50673, 8208]]
	<b>funds</b>	0.923120	0.004206	[[694055, 44], [12740, 27]]
	<b>mortgage</b>	0.927031	0.000000	[[702807, 0], [4059, 0]]
	<b>pension</b>	0.922354	0.005587	[[700441, 13], [6394, 18]]
	<b>loans</b>	0.854412	0.000000	[[705210, 0], [1656, 0]]
	<b>taxes</b>	0.859777	0.001202	[[668623, 21], [38199, 23]]
	<b>credit_card</b>	0.890683	0.006335	[[675711, 67], [30989, 99]]
	<b>securities</b>	0.914217	0.007435	[[689178, 57], [17565, 66]]
	<b>home_acct</b>	0.889242	0.000000	[[704156, 0], [2710, 0]]
	<b>pensions_2</b>	0.862506	0.001032	[[664248, 12], [42584, 22]]

	ROC AUC	F1 Score	Confusion Matrix
<b>direct_debt</b>	0.869315	0.046836	[[614537, 2327], [87788, 2214]]

```
In [17]: print("Validation Metrics Table:")
val_metrics_df
```

Validation Metrics Table:

Out[17]:

		ROC AUC	F1 Score	Confusion Matrix
Variation 1	savings_acct	0.448868	0.000000	[[200325, 0], [8, 0]]
	guarantees	0.117191	0.000000	[[200328, 3], [2, 0]]
	current_acct	0.756396	0.789899	[[39830, 38322], [17412, 104769]]
	derivada_acct	0.750282	0.000000	[[200289, 0], [44, 0]]
	payroll_acct	0.859517	0.007521	[[192120, 1999], [6183, 31]]
	junior_acct	0.997220	0.127197	[[199214, 0], [1043, 76]]
	mas_particular_acct	0.886404	0.000000	[[197523, 3], [2807, 0]]
	particular_acct	0.926977	0.245744	[[180445, 4825], [12277, 2786]]
	particular_plus_acct	0.858969	0.028845	[[194046, 1811], [4384, 92]]
	short_term_depo	0.912585	0.000000	[[200006, 3], [324, 0]]
	medium_term_depo	0.897696	0.000000	[[200152, 0], [181, 0]]
	long_term_depo	0.935919	0.293059	[[193029, 2643], [3407, 1254]]
	e_acct	0.876874	0.168256	[[187944, 2829], [8422, 1138]]
	funds	0.935718	0.058437	[[197675, 571], [2007, 80]]
	mortgage	0.946640	0.000000	[[199771, 41], [521, 0]]
	pension	0.932103	0.037847	[[198674, 780], [847, 32]]
	loans	0.858123	0.000000	[[199995, 0], [338, 0]]
	taxes	0.865688	0.030680	[[193529, 1867], [4831, 106]]
	credit_card	0.918365	0.040483	[[194089, 1880], [4235, 129]]
	securities	0.928463	0.085966	[[195880, 1401], [2852, 200]]
Variation 2	home_acct	0.899709	0.000000	[[200015, 0], [318, 0]]
	pensions_2	0.858688	0.009136	[[191616, 1993], [6684, 40]]
	direct_debt	0.870896	0.051411	[[182140, 2092], [15621, 480]]
	savings_acct	0.933433	0.000000	[[200325, 0], [8, 0]]
	guarantees	0.988746	0.000000	[[200331, 0], [2, 0]]
	current_acct	0.756892	0.789503	[[39858, 38294], [17517, 104664]]
	derivada_acct	0.907970	0.000000	[[200289, 0], [44, 0]]
	payroll_acct	0.860161	0.008951	[[192103, 2016], [6177, 37]]
	junior_acct	0.998574	0.827032	[[199092, 122], [244, 875]]
	mas_particular_acct	0.885058	0.000000	[[197474, 52], [2807, 0]]
	particular_acct	0.925391	0.253438	[[179507, 5763], [12041, 3022]]
	particular_plus_acct	0.861638	0.029767	[[193979, 1878], [4380, 96]]
	short_term_depo	0.945722	0.005831	[[199991, 18], [323, 1]]
	medium_term_depo	0.927063	0.037736	[[199969, 183], [174, 7]]

		ROC AUC	F1 Score	Confusion Matrix
	<b>long_term_depo</b>	0.938429	0.303536	[[192704, 2968], [3296, 1365]]
	<b>e_acct</b>	0.873791	0.171671	[[187851, 2922], [8388, 1172]]
	<b>funds</b>	0.938061	0.061080	[[196365, 1881], [1962, 125]]
	<b>mortgage</b>	0.948643	0.005950	[[197987, 1825], [514, 7]]
	<b>pension</b>	0.933830	0.031701	[[197539, 1915], [834, 45]]
	<b>loans</b>	0.887794	0.000000	[[199576, 419], [338, 0]]
	<b>taxes</b>	0.865983	0.031130	[[193439, 1957], [4828, 109]]
	<b>credit_card</b>	0.920199	0.044042	[[194071, 1898], [4223, 141]]
	<b>securities</b>	0.930085	0.096486	[[195539, 1742], [2809, 243]]
	<b>home_acct</b>	0.913726	0.009112	[[199896, 119], [316, 2]]
	<b>pensions_2</b>	0.859630	0.010488	[[191607, 2002], [6678, 46]]
	<b>direct_debt</b>	0.870889	0.060638	[[181975, 2257], [15527, 574]]
<b>Variation 3</b>	<b>savings_acct</b>	0.931511	0.000000	[[200325, 0], [8, 0]]
	<b>guarantees</b>	0.910590	0.000000	[[200315, 16], [2, 0]]
	<b>current_acct</b>	0.758703	0.790143	[[39445, 38707], [17107, 105074]]
	<b>derivada_acct</b>	0.897898	0.000000	[[200284, 5], [44, 0]]
	<b>payroll_acct</b>	0.864173	0.008957	[[192108, 2011], [6177, 37]]
	<b>junior_acct</b>	0.999552	0.840149	[[199085, 129], [215, 904]]
	<b>mas_particular_acct</b>	0.890196	0.000000	[[197523, 3], [2807, 0]]
	<b>particular_acct</b>	0.930324	0.257369	[[181266, 4004], [12247, 2816]]
	<b>particular_plus_acct</b>	0.866979	0.000000	[[195857, 0], [4476, 0]]
	<b>short_term_depo</b>	0.943757	0.012579	[[199701, 308], [320, 4]]
	<b>medium_term_depo</b>	0.928459	0.035794	[[199894, 258], [173, 8]]
	<b>long_term_depo</b>	0.940121	0.305180	[[192758, 2914], [3297, 1364]]
	<b>e_acct</b>	0.879096	0.170943	[[187857, 2916], [8394, 1166]]
	<b>funds</b>	0.942005	0.059712	[[197506, 740], [2000, 87]]
	<b>mortgage</b>	0.951792	0.007890	[[198818, 994], [515, 6]]
	<b>pension</b>	0.939714	0.036748	[[198570, 884], [846, 33]]
	<b>loans</b>	0.887644	0.000000	[[199658, 337], [338, 0]]
	<b>taxes</b>	0.873381	0.023401	[[194167, 1229], [4864, 73]]
	<b>credit_card</b>	0.923466	0.043750	[[194073, 1896], [4224, 140]]
	<b>securities</b>	0.931958	0.097025	[[195567, 1714], [2809, 243]]
	<b>home_acct</b>	0.917353	0.000000	[[200015, 0], [318, 0]]
	<b>pensions_2</b>	0.861479	0.010707	[[191601, 2008], [6677, 47]]

	ROC AUC	F1 Score	Confusion Matrix
<b>direct_debt</b>	0.870459	0.058622	[[181953, 2279], [15546, 555]]

Creating a summary table for all the variations and different datasets

```
In [18]: summary_data = []
for variation in metrics:
    for dataset in ['train', 'val']:
        avg_roc_auc = np.mean([metrics[variation][dataset][p]['ROC AUC'] for p in products])
        avg_f1 = np.mean([metrics[variation][dataset][p]['F1 Score'] for p in products])
        summary_data.append([variation, dataset, avg_roc_auc, avg_f1])

summary_df = pd.DataFrame(summary_data, columns=['Variation', 'Dataset', 'Avg ROC AUC', 'Avg F1 Score'])
print("Summary Table:")
print(summary_df.to_string(index=False))

best_variation = summary_df[summary_df['Dataset'] == 'val'].sort_values('Avg ROC AUC', ascending=False)
print(f"\nBest Model For This Week: {best_variation}")
```

Summary Table:

	Variation	Dataset	Avg ROC AUC	Avg F1 Score
Variation 1	train		0.827581	0.075498
Variation 1	val		0.836491	0.085847
Variation 2	train		0.887266	0.110663
Variation 2	val		0.907466	0.120787
Variation 3	train		0.888049	0.110734
Variation 3	val		0.906114	0.119955

Best Model For This Week: Variation 2

## Generate product recommendations

Here we want to visualize the product recommendations for each customer

```
In [19]: best_params = hyperparameter_variations[int(best_variation.split()[-1]) - 1]

# Train models and predict for each product for all customers using parameters from Variation 2
product_models = {}
for product in products:
    clf = LogisticRegression(**best_params)
    clf.fit(X_train, y_train[product])
    product_models[product] = clf

train_preds = {}
for product in products:
    proba = product_models[product].predict_proba(X_train)[ :, 1]
    for customer_id, prob in zip(train['customer_code'], proba):
        if customer_id not in train_preds:
            train_preds[customer_id] = []
        train_preds[customer_id].append((product, prob))
```

```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

```
n_iter_i = _check_optimize_result(
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.
```

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
`n_iter_i = _check_optimize_result(  
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.`

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

`n_iter_i = _check_optimize_result(  
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.`

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

`n_iter_i = _check_optimize_result(  
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.`

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

`n_iter_i = _check_optimize_result(  
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.`

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

`n_iter_i = _check_optimize_result(  
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.`

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

`n_iter_i = _check_optimize_result(  
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.`

Increase the number of iterations (max\_iter) or scale the data as shown in:

<https://scikit-learn.org/stable/modules/preprocessing.html>

Please also refer to the documentation for alternative solver options:

[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)

`n_iter_i = _check_optimize_result(  
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\linear_model\_logistic.py:460: ConvergenceWarning: lbfgs failed to converge (status=1):  
STOP: TOTAL NO. of ITERATIONS REACHED LIMIT.`



Increase the number of iterations (max\_iter) or scale the data as shown in:  
<https://scikit-learn.org/stable/modules/preprocessing.html>  
 Please also refer to the documentation for alternative solver options:  
[https://scikit-learn.org/stable/modules/linear\\_model.html#logistic-regression](https://scikit-learn.org/stable/modules/linear_model.html#logistic-regression)  
`n_iter_i = _check_optimize_result(`

We want to make sure that we are not recommending a product that the customer already own, so we will store the products that customers already have

```
In [20]: def get_active_products(customer_data):
         return set(product for product in products if customer_data[product] > 0)
```

We will sort the recommended products by the probability of a client getting it and we will get the top 7 recommendations

```
In [21]: for customer_id in train_preds:
         # Sorting by probability
         sorted_prods = sorted(train_preds[customer_id], key=lambda x: x[1], reverse=True)
         customer_data = train[train['customer_code'] == customer_id].iloc[0]
         # Filter out already active products
         active_products = get_active_products(customer_data)
         recommended_products = [prod for prod, _ in sorted_prods if prod not in active_products]

         # Get top 7
         train_preds[customer_id] = recommended_products[:7]

         # Example recommendations
         print("Recommendations:")
         for customer_id in list(train_preds.keys())[:5]:
             print(f"Customer {customer_id}: {train_preds[customer_id]}")
```

```
Recommendations:
Customer 1225385: ['current_acct', 'e_acct', 'particular_acct', 'taxes', 'particular_plus_acct', 'mas_particular_acct', 'direct_debt']
Customer 1358829: ['direct_debt', 'payroll_acct', 'pensions_2', 'e_acct', 'taxes', 'credit_card', 'long_term_depo']
Customer 1436539: ['direct_debt', 'payroll_acct', 'pensions_2', 'e_acct', 'short_term_depo', 'taxes', 'credit_card']
Customer 1448049: ['current_acct', 'e_acct', 'taxes', 'mas_particular_acct', 'direct_debt', 'home_acct', 'payroll_acct']
Customer 1396837: ['direct_debt', 'payroll_acct', 'pensions_2', 'mas_particular_acct', 'e_acct', 'junior_acct', 'long_term_depo']
```

Lastly, we want to identify which product has been recommended the most and least in the model

```
In [22]: product_rec_counts = {product: sum(1 for recs in train_preds.values() if product in recs)
         most_rec = max(product_rec_counts, key=product_rec_counts.get)
         least_rec = min(product_rec_counts, key=product_rec_counts.get)

         print(f"Most frequently recommended product: {most_rec} ({product_rec_counts[most_rec]} times)")
         print(f"Least frequently recommended product: {least_rec} ({product_rec_counts[least_rec]} times)")
```

```
Most frequently recommended product: taxes (652073 times)
Least frequently recommended product: savings_acct (2 times)
```