

# Week 11 — Explain the model, analyze risk, bias and ethical considerations

```
In [1]: import pandas as pd
import numpy as np
# import dask.dataframe as dd
import seaborn as sns
import matplotlib.pyplot as plt
from sklearn.model_selection import train_test_split
from datetime import datetime
from sklearn.metrics import roc_auc_score, f1_score, confusion_matrix
from sklearn.linear_model import LogisticRegression
from collections import defaultdict
from sklearn.preprocessing import MinMaxScaler
from joblib import Parallel, delayed
```

```
In [2]: pd.set_option('display.max_columns', None)

train = pd.read_csv('train_final.csv', low_memory=False)
validation = pd.read_csv('val_set_final.csv')
test = pd.read_csv('test_4_11.csv')
```

```
In [3]: train.head()
```

```
Out[3]:
```

	Unnamed: 0	date	customer_code	employee_index	country_spain	female	age	new_cust	sen
0	0	2016-04-28	1334092	N	1	0	0.234694	0	
1	1	2015-07-28	1024586	N	1	0	0.234694	0	
2	2	2016-04-28	856204	N	1	0	0.306122	0	
3	3	2015-08-28	295807	N	1	0	0.489796	0	
4	4	2016-03-28	942624	N	1	1	0.224490	0	

```
In [4]: validation.head()
```

Out[4]:

	Unnamed: 0	date	customer_code	employee_index	country_spain	female	age	first_contract_
0	0	2016-05-28	1212130	N	1	0	0.204082	2013-1
1	1	2015-07-28	84306	N	1	0	0.500000	1998-0
2	2	2015-07-28	883630	N	1	0	0.418367	2010-0
3	3	2016-05-28	1464700	N	1	1	0.183673	2015-0
4	4	2015-12-28	487783	N	1	1	0.418367	2004-1

In [5]: `test.head()`

Out[5]:

	Unnamed: 0	date	customer_code	employee_index	country_spain	female	age	new_cust	sen
0	0	2015-06-28	49335	N	1	0	0.734694	0	
1	1	2016-02-28	1174349	N	1	0	0.214286	0	
2	2	2015-07-28	1393286	N	1	0	0.244898	1	
3	3	2016-03-28	1454346	N	1	0	0.183673	0	
4	4	2016-02-28	1074431	N	1	0	0.234694	0	

Changing columns name and dropping columns so both datasets are the same

In [6]: `train = train.rename(columns={'country': 'country_spain'})`

In [7]: `train = train.drop(columns=['Unnamed: 0'])  
validation = validation.drop(columns=['Unnamed: 0'])  
drop = ['join_channel', 'province_name', 'employee_index', 'segment', 'total_products']  
train = train.drop(columns=drop)  
validation = validation.drop(columns=drop + ['payroll_acct.1', 'first_contract_date',  
  
test = test.drop(columns=['Unnamed: 0'])  
test = test.drop(columns=drop + ['payroll_acct.1'])`

## Reading into the data

Setting products we want to predict

```
In [8]: products = ['savings_acct', 'guarantees', 'current_acct',
                    'derivada_acct', 'payroll_acct', 'junior_acct', 'mas_particular_acct',
                    'particular_acct', 'particular_plus_acct', 'short_term_depo',
                    'medium_term_depo', 'long_term_depo', 'e_acct', 'funds', 'mortgage',
                    'pension', 'loans', 'taxes', 'credit_card', 'securities', 'home_acct',
                    'pensions_2', 'direct_debt']
```

## Transformation #1

Change #1: Instead of dropping these duplicates on customer column and use only the last instance we will keep those duplicates since it could capture some patterns such as if a client buys product x first, it will likely buy y product next.

We will create copies of the original train and test datasets so we don't change the original one.

```
In [9]: train_2 = train.copy()
        test_2 = test.copy()
```

## Pre-processing

## Transformation #2

For tranformation #2 we will add the date column as one of the features. For that, we will calculate the time since purchase using the month we are trying to predict on June 2016. For this transformation to make sense, we will also keep the first transformation, since the time line of purchase matters now, we will keep the duplicate clients' purchases instead of only keeping the last one

```
In [10]: train_2['date'] = pd.to_datetime(train_2['date'], format='%Y-%m-%d')

        train_2['date'] = train_2['date'].dt.to_period('M').dt.to_timestamp()

        # Setting our prediction date, June 28, 2016, as the reference date
        reference_date = pd.to_datetime("2016-06-28")

        # Calculate time since purchase
        train_2['months_since_purchase'] = (reference_date.year - train_2['date'].dt.year) * 12 +
                                           (reference_date.month - train_2['date'].dt.month)

        print(train_2[['date', 'months_since_purchase']])
```

	date	months_since_purchase
0	2016-04-01	2
1	2015-07-01	11
2	2016-04-01	2
3	2015-08-01	10
4	2016-03-01	3
...	...	...
5757281	2016-05-01	1
5757282	2015-08-01	10
5757283	2015-11-01	7
5757284	2016-05-01	1
5757285	2016-01-01	5

[5757286 rows x 2 columns]

```
In [11]: # Adding feature on test dataset
test_2['date'] = pd.to_datetime(test_2['date'], format='%Y-%m-%d')
test_2['date'] = test_2['date'].dt.to_period('M').dt.to_timestamp()

test_2['months_since_purchase'] = (reference_date.year - test_2['date'].dt.year) * 12
                                (reference_date.month - test_2['date'].dt.month)

print(test_2[['date', 'months_since_purchase']])
```

	date	months_since_purchase
0	2015-06-01	12
1	2016-02-01	4
2	2015-07-01	11
3	2016-03-01	3
4	2016-02-01	4
...	...	...
1236739	2016-02-01	4
1236740	2016-02-01	4
1236741	2015-08-01	10
1236742	2016-05-01	1
1236743	2016-04-01	2

[1236744 rows x 2 columns]

```
In [12]: X_train_2 = train_2.drop(['customer_code', 'date'] + products, axis=1)
y_train_2 = train_2[products]

X_test_2 = test_2.drop(['customer_code', 'date'] + products, axis=1)
y_test_2 = test_2[products]
```

## Training

```
In [13]: # Defining the best training parameter
params = {'C': 10, 'solver': 'liblinear', 'max_iter': 300}
```

Database with second transformation

```
In [14]: # Initialize dictionary for storing metrics
metrics = defaultdict(lambda: defaultdict(dict))

# Train and evaluate the model on the 'train_2' dataset
for product in products:
```

```

clf = LogisticRegression(**params)

# Train data and Labels for each product
y_train_2_product = y_train_2[product].values
y_test_2_product = y_test_2[product].values

# Train the model
clf.fit(X_train_2, y_train_2_product)

# Predictions
y_train_2_pred = clf.predict(X_train_2)
y_test_2_pred = clf.predict(X_test_2)
y_train_2_pred_proba = clf.predict_proba(X_train_2)[:, 1]
y_test_2_pred_proba = clf.predict_proba(X_test_2)[:, 1]

# Calculate metrics
metrics['train_2']['train'][product] = {
    'ROC AUC': roc_auc_score(y_train_2_product, y_train_2_pred_proba),
    'F1 Score': f1_score(y_train_2_product, y_train_2_pred),
    'Confusion Matrix': confusion_matrix(y_train_2_product, y_train_2_pred)
}

metrics['train_2']['test'][product] = {
    'ROC AUC': roc_auc_score(y_test_2_product, y_test_2_pred_proba),
    'F1 Score': f1_score(y_test_2_product, y_test_2_pred),
    'Confusion Matrix': confusion_matrix(y_test_2_product, y_test_2_pred)
}

```

```

In [15]: # Summarize the average metrics across all products
summary_data_2 = []
for dataset in ['train', 'test']:
    avg_roc_auc = np.mean([metrics['train_2'][dataset][p]['ROC AUC'] for p in products])
    avg_f1 = np.mean([metrics['train_2'][dataset][p]['F1 Score'] for p in products])
    summary_data_2.append(['train_2', dataset, avg_roc_auc, avg_f1])

# Create summary DataFrame
summary_df_2 = pd.DataFrame(summary_data_2, columns=['Dataset', 'Type', 'Avg ROC AUC', 'Avg F1 Score'])
print("\nEvaluated Model on Dataset: train_2")
print(summary_df_2.to_string(index=False))

```

```

Evaluated Model on Dataset: train_2
Dataset  Type  Avg ROC AUC  Avg F1 Score
train_2  train    0.876659    0.084174
train_2  test    0.874053    0.185602

```

## Feature Importance Analysis

```

In [ ]: # Dictionary to store feature importances
feature_importances = {}

# Iterate over each product
for product in products:
    clf = LogisticRegression(**params)
    clf.fit(X_train_2, y_train_2[product].values)

    feature_importances[product] = clf.coef_[0]

importances_df = pd.DataFrame(feature_importances, index=X_train_2.columns)

```

```
importances_df['Mean_Importance'] = importances_df.mean(axis=1)
sorted_importances = importances_df.sort_values(by='Mean_Importance', ascending=False)
```


In [27]: sorted\_importances

Out[27]:

	savings_acct	guarantees	current_acct	derivada_acct	payroll_acct	junior_acct
<b>active_cust</b>	0.830705	4.264000	1.414410	2.427595	6.842190	3.03010
<b>seniority_in_months</b>	4.603687	3.920184	0.148217	1.993498	0.956385	5.37899
<b>join_channel_encoded</b>	0.909934	-0.647531	-1.445911	0.368749	1.030798	2.54202
<b>province_name_encoded</b>	-0.805320	6.485177	-0.823748	-0.552544	0.626762	0.89281
<b>employee_index_encoded</b>	0.250569	0.528488	0.143572	0.657579	0.122628	0.48636
<b>months_since_purchase</b>	0.022000	0.026316	0.024692	0.010854	0.002302	0.03050
<b>income_to_age</b>	0.000004	-0.000041	-0.000020	0.000002	-0.000012	-0.00007
<b>income</b>	0.029702	0.014780	-0.000532	0.003247	-0.060334	-0.01263
<b>birth_spain</b>	0.615684	-3.797005	-0.192945	-0.808166	0.135772	-0.09866
<b>female</b>	-0.581870	-0.833165	0.029119	-1.419952	0.013041	0.00647
<b>cust_type</b>	-3.278677	-5.275618	0.531083	-0.524133	0.487662	-0.87524
<b>new_cust</b>	-2.920691	-3.626413	-0.168880	-1.249166	-0.418773	-0.66440
<b>country_spain</b>	-1.702266	-5.265817	0.772075	-2.432985	-3.752419	0.37180
<b>residency_spain</b>	-1.702266	-5.265817	0.772075	-2.432985	-3.752419	0.37180
<b>age</b>	-2.054414	-4.500836	-0.370727	1.399406	-2.178210	-69.70432

In [22]:

```
plt.figure(figsize=(10, 6))
sorted_importances['Mean_Importance'].plot(kind='bar', color='teal')
plt.title('Overall Feature Importance Across All Products')
plt.ylabel('Importance (Average Coefficient)')
plt.xlabel('Feature')
plt.xticks(rotation=45, ha='right')
plt.show()
```



## LIME example on variable 'savings\_acct'

Here we are trying to get deeper and understand which variables affect the variable 'savings\_acct' the most. For the purpose of this exercise we will just run the code on 'savings\_acct'

```
In [18]: # pip install lime
```

```
In [23]: from lime.lime_tabular import LimeTabularExplainer
from sklearn.preprocessing import StandardScaler
```

```
In [ ]: example_product = 'savings_acct'

#Standardizing features for LIME
scaler = StandardScaler()
X_train_scaled = scaler.fit_transform(X_train_2)
X_test_scaled = scaler.transform(X_test_2)

clf = LogisticRegression(**params)
clf.fit(X_train_scaled, y_train_2[example_product].values)

# Initializing LIME Tabular Explainer
explainer = LimeTabularExplainer(
    training_data=X_train_scaled,
    training_labels=y_train_2[example_product],
```

```



feature_names=X_train_2.columns,
class_names=['No', 'Yes'],
mode='classification'
)

# Choosing an instance to explain
instance_idx = 0
instance = X_test_scaled[instance_idx]

# Generating explanation
exp = explainer.explain_instance(
    data_row=instance,
    predict_fn=clf.predict_proba # Probability prediction function
)
exp.show_in_notebook(show_table=True) # For Jupyter Notebook
# exp.save_to_file('lime_explanation.html') # Save to HTML file

```

Prediction probabilities

No  1.00  
 Yes  0.00

No

Yes

seniority\_in\_months >...  
 0.00  
 cust\_type <= 0.05  
 0.00  
 new\_cust <= -0.18  
 0.00  
 birth\_spain <= -0.21  
 0.00  
 0.34 < join\_channel\_en...  
 0.00  
 age > 0.60  
 0.00  
 -0.89 < active\_cust <=...  
 0.00  
 province\_name\_encod...  
 0.00  
 female <= -0.91  
 0.00  
 employee\_index\_encod...  
 0.00

seniority_in_months	1.78
cust_type	0.05
new_cust	-0.18
birth_spain	-0.21
join_channel_encoded	1.21
age	1.95
active_cust	1.12
province_name_encoded	-0.88
female	-0.91
employee index encoded	-0.02

Select 5 predictions at random, explain how the model generated those predictions (which features matter more than others), which features need to change and by how much to move the output in a significant way (e.g., to flip the prediction from one class to another)



```
In [24]: import lime
import lime.lime_tabular
import random
```

```
In [25]: # Randomly selecting 5 samples
random_indices = random.sample(range(X_test_2.shape[0]), 5)
selected_samples = X_test_2.iloc[random_indices]
selected_labels = y_test_2[product].iloc[random_indices]

# LIME explainer
explainer = lime.lime_tabular.LimeTabularExplainer(
    X_train_2.values,
    feature_names=X_train_2.columns,
    class_names=[f"Not {product}", product],
    verbose=True,
    mode="classification"
)

# Explain each sample
for i, idx in enumerate(random_indices):
    print(f"\nExplaining Prediction {i+1} (Row {idx}):")
    sample = X_test_2.iloc[idx].values

    # Probability for the sample
    prediction = clf.predict_proba([sample])[0]
    predicted_class = np.argmax(prediction)
    print(f"Predicted Class: {predicted_class} (Probability: {prediction[predicted_class]})")

    # Generate explanation
    exp = explainer.explain_instance(sample, clf.predict_proba, num_features=10)
    exp.show_in_notebook(show_table=True)

    # Print explanation in text format
    explanation = exp.as_list()
    print("Feature Contributions:")
    for feature, contribution in explanation:
        print(f"{feature}: {contribution:.4f}")


    # Identify feature changes to flip prediction
    important_feature, contribution = explanation[0]
    print(f"\nTo flip the prediction, try adjusting '{important_feature}' by a significant amount")

# Visualize the explanation for one of the samples
exp.show_in_notebook(show_table=True)
```

```
Explaining Prediction 1 (Row 950916):
Predicted Class: 1 (Probability: 0.87)
Intercept -0.039240922130545156
Prediction_local [0.15935189]
Right: 0.8716594547837396
```

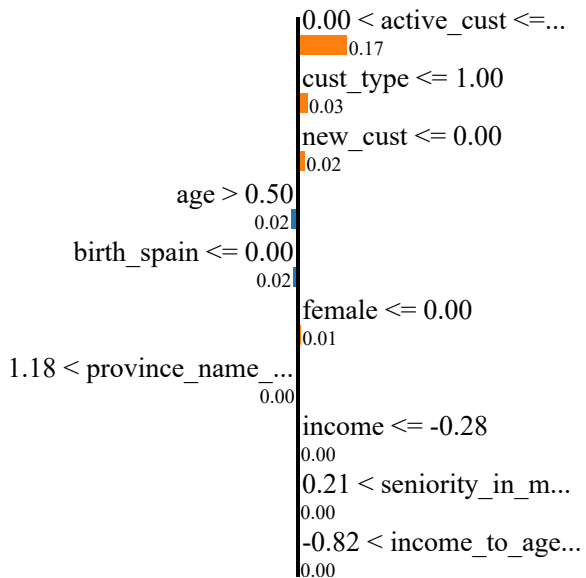
```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
  warnings.warn(
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
  warnings.warn(
```

## Prediction probabilities



Not direct\_debt

direct\_debt



## Feature Value

active_cust	1.00
cust_type	1.00
new_cust	0.00
age	0.66
birth_spain	0.00
female	0.00
province_name_encoded	1.24
income	-0.46
seniority_in_months	0.54

## Feature Contributions:

0.00 < active\_cust <= 1.00: 0.1687

cust\_type <= 1.00: 0.0320

new\_cust <= 0.00: 0.0229

age > 0.50: -0.0219

birth\_spain <= 0.00: -0.0151

female <= 0.00: 0.0084

1.18 < province\_name\_encoded <= 1.28: -0.0037

income <= -0.28: 0.0034

0.21 < seniority\_in\_months <= 0.54: 0.0023

-0.82 < income\_to\_age <= -0.35: 0.0016

To flip the prediction, try adjusting '0.00 < active\_cust <= 1.00' by a significant amount.

## Explaining Prediction 2 (Row 3239):

Predicted Class: 0 (Probability: 0.96)


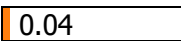
Intercept 0.21402389014112455

Prediction\_local [0.03316361]

Right: 0.037688068101520936

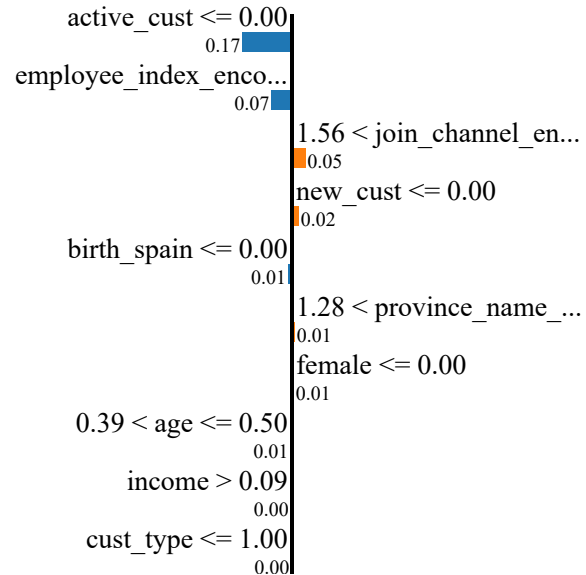
```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
  warnings.warn(
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
  warnings.warn(
```

Prediction probabilities

Not direct\_debt  0.96  
 direct\_debt  0.04

Not direct\_debt

direct\_debt



Feature Value

active_cust	0.00
employee_index_encoded	1.41
join_channel_encoded	1.81
new_cust	0.00
birth_spain	0.00
province_name_encoded	1.75
female	0.00
age	0.50
income	0.45

## Feature Contributions:

```
active_cust <= 0.00: -0.1708
employee_index_encoded <= 1.41: -0.0710
1.56 < join_channel_encoded <= 1.94: 0.0472
new_cust <= 0.00: 0.0232
birth_spain <= 0.00: -0.0136
1.28 < province_name_encoded <= 1.75: 0.0093
female <= 0.00: 0.0069
0.39 < age <= 0.50: -0.0056
income > 0.09: -0.0043
cust_type <= 1.00: -0.0021
```

To flip the prediction, try adjusting 'active\_cust <= 0.00' by a significant amount.

## Explaining Prediction 3 (Row 458596):

Predicted Class: 0 (Probability: 0.99)

Intercept 0.30015283526118897

Prediction\_local [-0.02265368]

Right: 0.014329623063726801

```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
```

```
warnings.warn(
```

```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
```

```
warnings.warn(
```

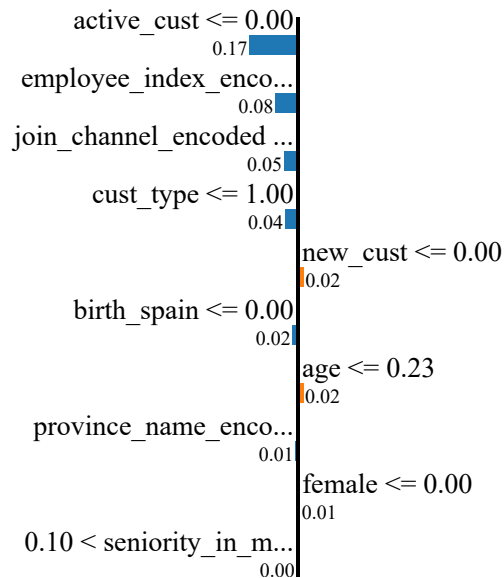
## Prediction probabilities

Not direct\_debt 0.99

direct\_debt 0.01

Not direct\_debt

direct\_debt



## Feature Value

active_cust	0.00
employee_index_encoded	1.41
join_channel_encoded	0.89
cust_type	1.00
new_cust	0.00
birth_spain	0.00
age	0.20
province_name_encoded	1.18
female	0.00

## Feature Contributions:

active\_cust <= 0.00: -0.1668

employee\_index\_encoded <= 1.41: -0.0780

join\_channel\_encoded <= 0.89: -0.0475

cust\_type <= 1.00: -0.0443

new\_cust <= 0.00: 0.0195

birth\_spain <= 0.00: -0.0178

age <= 0.23: 0.0173

province\_name\_encoded <= 1.18: -0.0074

female <= 0.00: 0.0067

0.10 < seniority\_in\_months <= 0.21: -0.0044

To flip the prediction, try adjusting 'active\_cust <= 0.00' by a significant amount.

## Explaining Prediction 4 (Row 1115023):

```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
  warnings.warn(
```

Predicted Class: 1 (Probability: 0.95)  
 Intercept -0.03681746983404906  
 Prediction\_local [0.22245744]  
 Right: 0.9472943668943028

c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names  
 warnings.warn(

Prediction probabilities

Not direct\_debt 0.05  
 direct\_debt 0.95

Not direct\_debt

direct\_debt

0.00 < active\_cust <= ... 0.17  
 1.56 < join\_channel\_en... 0.05  
 new\_cust <= 0.00 0.03  
 birth\_spain <= 0.00 0.01  
 1.28 < province\_name\_... 0.01  
 female <= 0.00 0.01  
 0.23 < age <= 0.39 0.01  
 employee\_index\_enco... 0.01  
 seniority\_in\_months ... 0.00  
 income\_to\_age > 0.23 0.00

active_cust	1.00
join_channel_encoded	1.94
new_cust	0.00
birth_spain	0.00
province_name_encoded	1.75
female	0.00
age	0.36
employee_index_encoded	1.41
seniority_in_months	0.05
income to age	0.51

Feature Contributions:

```
0.00 < active_cust <= 1.00: 0.1695
1.56 < join_channel_encoded <= 1.94: 0.0546
new_cust <= 0.00: 0.0291
birth_spain <= 0.00: -0.0150
1.28 < province_name_encoded <= 1.75: 0.0127
female <= 0.00: 0.0079
0.23 < age <= 0.39: 0.0071
employee_index_encoded <= 1.41: -0.0055
seniority_in_months <= 0.10: -0.0045
income_to_age > 0.23: 0.0035
```

To flip the prediction, try adjusting '0.00 < active\_cust <= 1.00' by a significant amount.

Explaining Prediction 5 (Row 557414):

Predicted Class: 0 (Probability: 0.99)

```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
  warnings.warn(
```

Intercept 0.1629545436552834

Prediction\_local [-0.02360619]

Right: 0.014531659456663631

```
c:\Users\MARIA\AppData\Local\Programs\Python\Python311\Lib\site-packages\sklearn\base.py:464: UserWarning: X does not have valid feature names, but LogisticRegression was fitted with feature names
  warnings.warn(
```

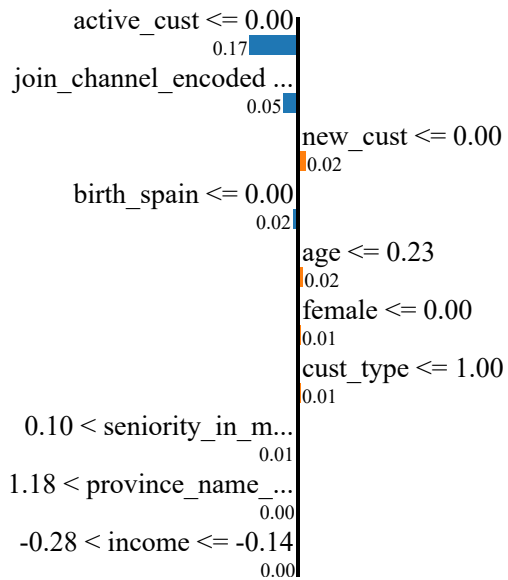
## Prediction probabilities

Not direct\_debt 0.99

direct\_debt 0.01

Not direct\_debt

direct\_debt



## Feature Value

active_cust	0.00
join_channel_encoded	0.89
new_cust	0.00
birth_spain	0.00
age	0.23
female	0.00
cust_type	1.00
seniority_in_months	0.21
province_name_encoded	1.28

## Feature Contributions:

active\_cust <= 0.00: -0.1665

join\_channel\_encoded <= 0.89: -0.0483

new\_cust <= 0.00: 0.0249

birth\_spain <= 0.00: -0.0164

age <= 0.23: 0.0161

female <= 0.00: 0.0079

cust\_type <= 1.00: 0.0074

0.10 < seniority\_in\_months <= 0.21: -0.0063


1.18 < province\_name\_encoded <= 1.28: -0.0042

-0.28 < income <= -0.14: -0.0012

To flip the prediction, try adjusting 'active\_cust <= 0.00' by a significant amount.

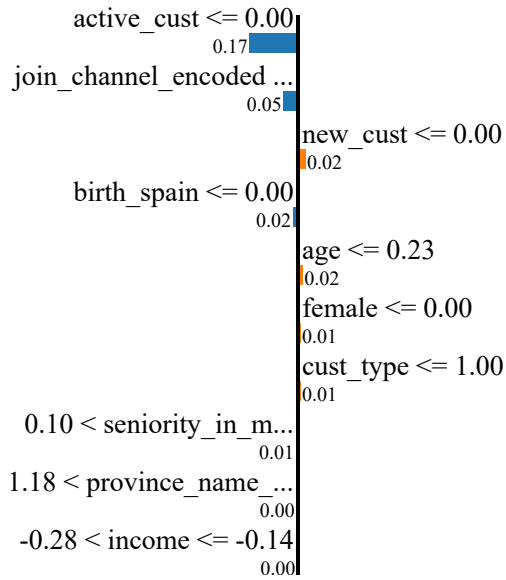


Prediction probabilities



Not direct\_debt

direct\_debt



Feature Value

active_cust	0.00
join_channel_encoded	0.89
new_cust	0.00
birth_spain	0.00
age	0.23
female	0.00
cust_type	1.00
seniority_in_months	0.21
province_name_encoded	1.28